

EXPERIMENT NO. 8

Date: _____

TITLE

Ridge and Lasso Regression on Diabetes Dataset

AIM

To fit ordinary, Ridge and Lasso models on diabetes data, pick penalties by 5-fold cross-validation, and compare test MSE and R^2 .

OBJECTIVE

To implement Ridge and Lasso regression on the Diabetes dataset, tune penalties by cross-validation, and compare with standard linear regression using MSE and R-squared.

THEORY

Standard linear regression (normal equation):

$$\theta = (X^T X)^{-1} X^T y$$

Ridge objective and closed form (λ unpenalized intercept via centering):

$$J = \|y - Xw\|^2 + \lambda \|w\|^2 \quad w = (X^T X + \lambda I)^{-1} X^T y$$

Lasso objective and coordinate-descent update (soft thresholding):

$$J = \|y - Xw\|^2 + \lambda \|w\|_1 \quad w_j = \frac{S(\rho_j, \lambda)}{z_j}, \quad S(\rho, \lambda) = \text{sign}(\rho) \max(|\rho| - \lambda, 0)$$

Cross-validation picks the λ with lowest mean validation MSE; report test MSE and R^2 .

PROCEDURE

1. Load diabetes data (442×10); 80/20 split (seed 42); standardize X , center y (intercept unpenalized).
2. Fit ordinary least squares by normal equation; record test MSE and R^2 .
3. Fit Ridge in closed form over a λ grid; pick best λ by 5-fold cross-validation.
4. Fit Lasso by coordinate descent over an α grid; pick best α by 5-fold cross-validation.
5. Compare test MSE, R^2 and sparsity; plot CV curves.

DATASET DESCRIPTION

Dataset Name: Diabetes (diabetes.tab.txt)

Source: 442 patients; 10 baseline variables (age, sex, bmi, bp, s1–s6)

Number of Samples: 353 train / 89 test Number of Features: 10

Target / Class: Disease progression (continuous)

Description: 80/20 split, seed 42; X standardized, y centered.

PROGRAM

```

import numpy as np
import matplotlib
matplotlib.use("Agg")
import matplotlib.pyplot as plt
CSV = "diabetes.tab.txt"
data = np.loadtxt(CSV, delimiter="\t", skiprows=1)
X_all, y_all = data[:, :10], data[:, 10]
print("Dataset:", X_all.shape)
rng = np.random.default_rng(42)
idx = np.arange(len(X_all))
rng.shuffle(idx)
ntr = int(0.8 * len(idx))
tri, tei = idx[:ntr], idx[ntr:]
Xtr_raw, Xte_raw, ytr, yte = X_all[tri], X_all[tei], y_all[tri], y_all[tei]
print("Train:", len(tri), " Test:", len(tei))
mu, sg = Xtr_raw.mean(0), Xtr_raw.std(0); yb = ytr.mean()
Xtr, Xte = (Xtr_raw - mu) / sg, (Xte_raw - mu) / sg; ytrc = ytr - yb
def ridge_fit(X, y, lam):
    p = X.shape[1]
    return np.linalg.inv(X.T @ X + lam * np.eye(p)) @ X.T @ y
def lasso_fit(X, y, lam, iters=300):
    n, p = X.shape
    w = np.zeros(p)
    for _ in range(iters):
        for j in range(p):
            rho = X[:, j] @ (y - X @ w + w[j] * X[:, j]) / n
            z = (X[:, j] ** 2).mean()
            w[j] = np.sign(rho) * max(abs(rho) - lam, 0.0) / z
    return w
def mse(a, b):
    return float(((a - b) ** 2).mean())
def r2(ys, p):
    return float(1 - ((ys - p) ** 2).sum() / ((ys - ys.mean()) ** 2).sum())
w_ols = np.linalg.inv(Xtr.T @ Xtr) @ Xtr.T @ ytrc
p_ols = Xte @ w_ols + yb
print(f"OLS: MSE={mse(yte, p_ols):.4f} R2={r2(yte, p_ols):.4f}")
folds = np.array_split(rng.permutation(ntr), 5)
def cv_score(lam, lasso):
    e = 0.0
    for k in range(5):
        va = folds[k]
        tr = np.concatenate([folds[j] for j in range(5) if j != k])
        w = (lasso_fit if lasso else ridge_fit)(Xtr[tr], ytrc[tr], lam)
        e += ((ytr[va] - (Xtr[va] @ w + yb)) ** 2).mean()
    return e / 5
rlams = [0.1, 1.0, 10.0, 100.0]; llams = [0.1, 0.5, 1.0, 5.0]
rms = [cv_score(l, False) for l in rlams]; lms = [cv_score(l, True) for l in llams]
bl, ba = rlams[int(np.argmin(rms))], llams[int(np.argmin(lms))]
print(f"Best: Ridge lam={bl}, Lasso alpha={ba}")
w_r = ridge_fit(Xtr, ytrc, bl); w_l = lasso_fit(Xtr, ytrc, ba)
p_r, p_l = Xte @ w_r + yb, Xte @ w_l + yb
print(f"Ridge: MSE={mse(yte, p_r):.4f} R2={r2(yte, p_r):.4f}")
print(f"Lasso: MSE={mse(yte, p_l):.4f} R2={r2(yte, p_l):.4f}")
print("Nonzero Ridge/Lasso:", (w_r != 0).sum(), (w_l != 0).sum())
plt.figure(figsize=(7, 4.5))
plt.semilogx(rlams, rms, marker="o", linewidth=2, label="Ridge CV")
plt.semilogx(llams, lms, marker="s", linewidth=2, label="Lasso CV")
plt.xlabel("Lambda"); plt.ylabel("5-fold CV MSE"); plt.grid(True)
plt.title("Cross-Validation: Ridge vs Lasso"); plt.legend(); plt.tight_layout()
plt.savefig("exp8_cv.png", dpi=150)

```

OUTPUT

Dataset: (442, 10)
Train: 353 Test: 89
OLS: MSE=2594.6297 R2=0.5927
Best: Ridge lam=1.0, Lasso alpha=0.5
Ridge: MSE=2596.3521 R2=0.5924
Lasso: MSE=2623.3691 R2=0.5882
Nonzero Ridge/Lasso: 10 8

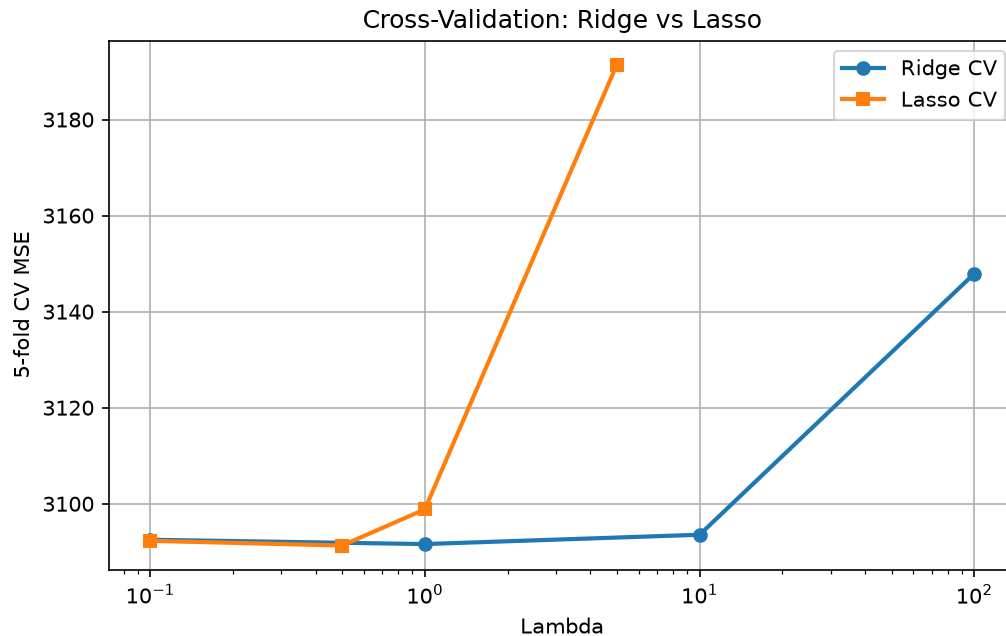


Fig. 1: 5-fold CV MSE vs penalty (best: Ridge $\lambda = 1$, Lasso $\alpha = 0.5$).

RESULT

The above program was implemented and executed successfully using Python, and the corresponding output was obtained and verified. The results of the **Ridge and Lasso regression** were analyzed and found to be satisfactory, thereby achieving the desired objective of the experiment.

Faculty Signature: _____

Date: _____